

Combinatorial optimization on quantum computers

Ashley Montanaro

Phasecraft Ltd
School of Mathematics, University of Bristol

Combinatorial optimization

Combinatorial optimization problems are characterised by needing to search over **exponentially many possible solutions**.

For example:

- **Colouring a graph** with the minimal number of colours such that no adjacent vertices share a colour;
- Finding the **lowest-cost route** that visits all of a set of cities;
- Determining if a **system of linear equations** over integers $\{0, 1\}$ has a solution.

Quantum computers can sometimes achieve a speedup in solving optimization problems over our best classical algorithms.

However, the speedup (when provable) is usually at most quadratic, as opposed to exponential (e.g. running time $2^n \rightarrow 2^{n/2}$)

Today's talk

Today I will discuss some quantum algorithms for solving optimization and constraint satisfaction problems (CSPs):


- Quantum speedup of **backtracking** and **branch-and-bound** algorithms [AM '15, AM '19] (with a large-scale, fault-tolerant quantum computer)
- Quantum speedup of **boolean satisfiability** [Boulebnane and AM '23] (possibly with a near-term quantum computer)

Solving hard constraint satisfaction problems with structure

Often we can achieve a better complexity than unstructured search or optimization by using the **structure** of the problem we need to solve.

One of the most prominent techniques for this is **backtracking** (“trial and error”).

We can illustrate backtracking with graph k -colouring:



An **NP-complete** problem with a huge number of direct applications, including register allocation; scheduling; frequency assignment problems; ...


Colouring by backtracking




Colouring by backtracking




Colouring by backtracking




Colouring by backtracking




Colouring by backtracking



Colouring by backtracking




Colouring by backtracking



Search in a tree

Imagine we want to find a “marked” vertex in a tree where we only have local knowledge, starting from the root.




If the tree has T vertices, this requires $\sim T$ time classically in the worst case.

Quantum search in a tree

Theorem [Belovs '13]



There is a quantum algorithm that can detect existence of a marked vertex in a tree with T vertices and depth d , using $O(\sqrt{Td})$ queries.



The algorithm is based on a **quantum walk** in the tree.

From quantum search in trees to backtracking

- A backtracking algorithm solving a problem with n variables explores a tree of size T and depth n .
- The quantum walk algorithm for search in trees can be applied, yielding:

Theorem (informal) [AM '18]


There is a corresponding quantum algorithm which finds a solution, or outputs that one does not exist, in time $O(\sqrt{T} \text{poly}(n))$, with 1% probability of error.

- We normally think of T as exponential in n ; in this case, the speedup is **near-quadratic**.
- Subsequent improvements to this algorithm: [Ambainis and Kokainis '17], [Jarret and Wan '18]

Branch-and-bound algorithms

The backtracking approach can be generalised to solve **optimization problems** via a technique known as branch-and-bound, which is applicable whenever we have:

- A **branching** procedure that splits a set of potential solutions into subsets;
- A **bounding** procedure that returns a lower bound on the cost of any solution in a subset.



Quantum speedup of branch-and-bound

Theorem (informal) [AM'19]

Assume there is a classical algorithm that solves an optimization problem using the branch and bound procedures T times. Then there is a quantum algorithm that solves the same problem using these procedures $O(\sqrt{T})$ times (up to lower-order terms).

- The quantum algorithm is based on the use of the backtracking algorithm as a subroutine.
- It can be applied to find ground states of the **Ising model** (aka **Max-Cut**):

$$\min_{z \in \{\pm 1\}^n} \sum_{i < j} a_{ij} z_i z_j$$

- e.g. Sherrington-Kirkpatrick model $a_{ij} \sim N(0, 1)$: runtime $O(2^{0.226n})$ or better, beating Grover search $O(2^{0.5n})$.

Other developments in quantum backtracking

Applications:

- lattice-based cryptography, e.g. [Alkim et al '16, del Pino et al '16]
- Travelling Salesman Problem [Moylett et al '17]
- exact satisfiability [Mandrà et al '16]
- constraint programming [Booth et al '21]

Experimental implementation (in simulation) for 2-colouring a graph with 4 vertices [Martiel and Remaud '19]

Do these theoretical speedups translate into real-world speedups?

Imagine we want to solve a problem within a runtime of at most 1 day.

- In an optimistic hardware parameter regime, we could see speedup factors of $> 10^4$ (compared with a standard desktop PC) [Campbell et al '19]
- This speedup gets substantially smaller when considering parameters corresponding to quantum hardware available today.
- If we additionally take into account the cost of classical error-correction processing, this speedup disappears.
- The number of physical qubits used is very large (e.g. $> 10^{12}$), almost all of which are used for fault-tolerance.
- This strongly motivates the design of improved fault-tolerance techniques!

Boolean satisfiability

Are there hard constraint satisfaction problems which we may be able to solve using quantum computers in the **near term**?

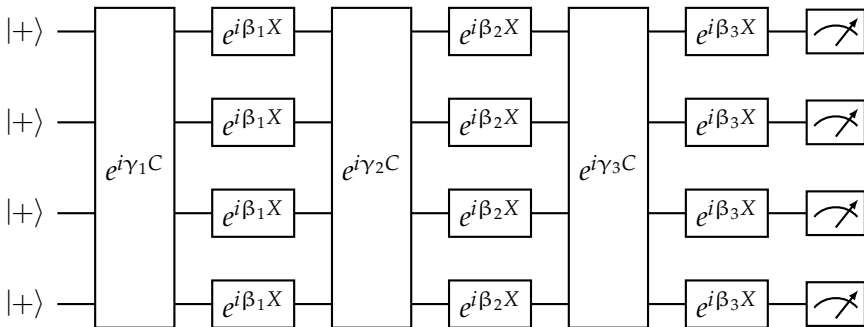
Consider the **boolean satisfiability** problem: finding a solution to a set of constraints on boolean variables x_1, \dots, x_n of the form

$$(x_1 \vee x_2 \vee x_3) \wedge (\neg x_1 \vee x_4 \vee x_7) \wedge (x_2 \vee \neg x_5 \vee \neg x_7)$$

Here: the special case of the **random k -SAT** problem where each clause contains k variables and is chosen **at random**.

The quantum approximate optimization algorithm (QAOA)

We will apply a quantum algorithm known as QAOA [Hogg '00, Farhi et al '14] for minimising a cost function $C : \{0, 1\}^n \rightarrow \mathbb{Z}$.



The quantum approximate optimization algorithm (QAOA)

Mathematically, QAOA looks like this:

- Alternate between evolution under (diagonal) **cost** Hamiltonian and **transverse field** Hamiltonian.

$$|\Psi_{\text{QAOA}}(\boldsymbol{\beta}, \boldsymbol{\gamma})\rangle := e^{-i\beta_p H_B} e^{-i\gamma_p H_C} \dots e^{-i\beta_1 H_B} e^{-i\gamma_1 H_C} |+\rangle^{\otimes n},$$

$$H_C := \sum_{z \in \{0,1\}^n} C(z) |z\rangle \langle z|,$$

$$H_B := \sum_{1 \leq j \leq n} X_j,$$

where C is an integer-valued numerical function.

- After evolution, **measure state in computational basis** to obtain candidate solution.
- **Variational quantum algorithm**: parameters $\boldsymbol{\beta}, \boldsymbol{\gamma}$ (“QAOA angles”) need to be determined.

Applying QAOA to k -SAT

$$(x_1 \vee x_2 \vee x_3) \wedge (\neg x_1 \vee x_4 \vee x_7) \wedge (x_2 \vee \neg x_5 \vee \neg x_7)$$

- Cost function $C(z)$ = **number of violated clauses**.

Problem instance: $\sigma = (\sigma_0, \dots, \sigma_{m-1})$

$$\text{Cost function: } C_{\sigma}(z) := \sum_{j \in [m]} \mathbf{1} [z \text{ violates } \sigma_j]$$

$$\text{Diagonal Hamiltonian: } H[\sigma] := \sum_{z \in \{0,1\}^n} C_{\sigma}(z) |z\rangle \langle z|$$

- We wish to **maximize the success probability** $p_{\text{success}}(\sigma, \beta, \gamma)$ (not the expected number of satisfied constraints!):

$$\max_{\beta, \gamma} \langle \Psi_{\text{QAOA}}(\sigma, \beta, \gamma) | \Pi_{\ker H[\sigma]} | \Psi_{\text{QAOA}}(\sigma, \beta, \gamma) \rangle$$

- We do this by optimizing the parameters over random small instances.

Analytic estimate for the success probability

For **random instances** of k -SAT with an expected number of clauses $m = rn$, we show:

Exponential scaling of expected success probability

Let $k = 2^q \geq 2$ be an integer, $p \geq 1$ an integer and $\beta, \gamma \in \mathbf{R}^p$.
For γ sufficiently small (independent of instance size n),

$$\lim_{n \rightarrow \infty} \frac{1}{n} \log \mathbf{E}_{\sigma \sim k\text{SAT}(n, r, 2^q)} [p_{\text{success}}(\sigma, \beta, \gamma)]$$

exists and can be computed from the fixed point of a function of 2^{2p+1} complex variables.

The scaling exponent of the expected success probability calculated by the procedure can be used as a proxy to **optimize the success probability** over β, γ .

Analytic vs. empirical scaling exponents

How well do analytic predictions match empirical results?





Figure: Analytic and empirical exponents for fixed $k = 8$ and varying p . Empirical exponents were fitted from empirical average success probabilities (solid lines). Analytic fits (dashed lines) were represented assuming a success probability 1 “at $n = 0$ ”.

QAOA vs. classical algorithms

- We benchmarked QAOA against many classical SAT solvers; **WalkSATlm** found to perform best.
- Method: similar to benchmarks above but consider **median** running time, not average-instance success probability.



Solver	Fit
WalkSAT QAOA	$-3.232 + 0.295n$
QAOA ($p = 14$)	$-1.064 + 0.326n$
QAOA ($p = 60$)	$-2.842 + 0.302n$
walksatlm	$-0.309 + 0.325n$
maplesat	$1.531 + 0.461n$
glucose4	$2.998 + 0.498n$

Figure: Scaling behaviour of median running times of selected classical and quantum algorithms for 8-SAT. WalkSAT QAOA uses $p = 60$.

QAOA vs. WalkSATlm

How deep a quantum algorithm will we need to outperform classical algorithms?




Figure: Running times of QAOA compared with WalkSATlm for random 8-SAT. Blue dashed line is fitting based on all p , blue solid line is using $p \leq 10$.

Conclusions

We might be able to achieve a fairly significant quantum speedup for common and practically relevant combinatorial optimization problems...
...but we're not quite there yet! More work remains to be done on algorithm design and applications.

Advert: We are hiring at Phasecraft for quantum algorithms scientists to work on **quantum optimization algorithms**.

Further reading:


- **Backtracking**: arXiv:1509.02374, arXiv:1906.10375
- **k-SAT**: arXiv:2208.06909 (with Sami Boulebnane)
- **Survey on quantum optimization algorithms**: arXiv:2312.02279 (Abbas et al)

Thanks!

Quantum walk in a tree

For each vertex x with c children, define

$$|\psi_x\rangle = \frac{1}{\sqrt{c+1}} \left(|x\rangle + \sum_{\text{children } y} |y\rangle \right), \quad D_x = \begin{cases} I - 2|\psi_x\rangle\langle\psi_x| & [x \text{ unmarked}] \\ I & [x \text{ marked}]. \end{cases}$$




e.g. $|\psi_2\rangle = \frac{1}{\sqrt{3}} (|2\rangle + |4\rangle + |5\rangle), \quad D_2 = \frac{1}{3} \begin{pmatrix} 1 & -2 & -2 \\ -2 & 1 & -2 \\ -2 & -2 & 1 \end{pmatrix}$

Quantum walk in a tree

Let A and B be the sets of vertices an even and odd distance from the root, respectively.


Then a step of the walk is applying the unitary operator $W = D_B D_A$, where $D_A = \bigoplus_{x \in A} D_x$ and $D_B = \bigoplus_{x \in B} D_x$.



Quantum walk in a tree

Let A and B be the sets of vertices an even and odd distance from the root, respectively.


Then a step of the walk is applying the unitary operator $W = D_B D_A$, where $D_A = \bigoplus_{x \in A} D_x$ and $D_B = \bigoplus_{x \in B} D_x$.



Quantum walk in a tree

Let A and B be the sets of vertices an even and odd distance from the root, respectively.


Then a step of the walk is applying the unitary operator $W = D_B D_A$, where $D_A = \bigoplus_{x \in A} D_x$ and $D_B = \bigoplus_{x \in B} D_x$.



Quantum walk in a tree

Let A and B be the sets of vertices an even and odd distance from the root, respectively.

Then a step of the walk is applying the unitary operator $W = D_B D_A$, where $D_A = \bigoplus_{x \in A} D_x$ and $D_B = \bigoplus_{x \in B} D_x$.



Analysing the quantum walk

Claims

- 1 ≥ 1 marked vertex: W has an eigenvector w /eigenvalue 1.
- 2 No marked vertex: W has no eigenvector w /eigenvalue 1.

Quantum computers can estimate eigenvalues via a procedure called **phase estimation**, allowing these two cases to be distinguished.

Small print: We actually need to apply this to a similar unitary operator W' , such that the root $|r\rangle$ is close to an eigenvector of W' .

W' has an approximate “phase gap” of $\sim 1/\sqrt{Td}$, allowing the two cases to be distinguished with $O(\sqrt{Td})$ uses of W' .

Cost model

We work out the runtime and space usage of quantum algorithms based on the use of the **surface code** [Fowler et al '12] for quantum error-correction.

We then convert this to real-world runtimes based on various regimes corresponding to different parameters for quantum-computing hardware:

Parameter	Realistic	Plausible	Optimistic
Measurement time	50ns	5ns	0.5ns
2-qubit gate time	30ns	3ns	0.3ns
Gate error rate	10^{-3}	10^{-4}	10^{-5}

“Realistic” is (approximately!) achievable today; other two columns represent order-of-magnitude improvements.

Summary of results

	Realistic	Plausible	Optimistic
Max n	113	128	144
T-depth	1.70×10^{12}	1.53×10^{13}	1.62×10^{14}
T/Toffoli count	8.24×10^{17}	9.94×10^{18}	1.24×10^{20}
Factory qubits	6.29×10^{13}	9.26×10^{12}	3.59×10^{12}
Speedup factor	7.25×10^0	5.17×10^2	4.16×10^4

Table: Likely speedup factors for graph colouring via backtracking achievable in different regimes.

Complexity estimates for other algorithms and CSPs were obtained by [\[Sanders et al '20\]](#).